

Assignment 5 - Build and Deploy a Keyword Spotting Model using Edge Impulse

Q. Does the model perform as accurately as expected on your smartphone? List a few methods to improve the model's accuracy.

The model's performance on the smartphone aligns with expectations but slightly lower compared to the BLE module.

Methods to Improve Model Accuracy

- **Enhance Data Quality and Diversity:**
 - Collect a larger and more diverse dataset to cover various accents, noise levels, and environments.
 - Add data augmentation techniques like noise injection, pitch shifting, or time-stretching to improve the model's generalization.
- **Optimize Preprocessing Steps:**
 - Experiment with different feature extraction techniques, such as using higher-resolution MFCCs or exploring other audio features like spectrograms.
- **Improve Model Architecture:**
 - Try more advanced lightweight architectures like MobileNet or TinyML-specific models.
 - Add regularization techniques (e.g., dropout) to prevent overfitting during training.
- **Fine-Tune Hyperparameters:**
 - Experiment with learning rates, batch sizes, and epochs to improve the model's training and testing performance.
- **Post-Training Optimization:**
 - Quantize the model to FP16 or INT8 to make it suitable for deployment without significantly impacting accuracy.
 - Apply techniques like pruning to remove less critical model parameters while maintaining performance.

Q. When building a model for resource-limited hardware, how do you balance fast inference times with acceptable model accuracy? What trade-offs did you encounter?

When building models for resource-limited hardware like the BLE module, there is often a trade-off between achieving faster inference times and maintaining acceptable accuracy. To strike the right balance:

Key Considerations

- **Model Size and Complexity:**
 - Smaller models are faster but can compromise accuracy. Selecting architectures optimized for TinyML, such as MobileNet or SqueezeNet, can help strike a balance.
- **Precision Reduction:**
 - Quantization techniques (e.g., INT8 quantization) can significantly improve inference speed and reduce memory usage but may slightly lower accuracy.
- **Feature Simplification:**
 - Reducing the dimensionality of features (e.g., using fewer MFCC coefficients) can speed up inference but might lead to less precise classifications.

Trade-Offs Encountered

- **Reduced Precision vs. Accuracy:**
 - Lowering the model precision to reduce memory usage and speed up inference might slightly degrade the accuracy.
- **Simplified Features vs. Robustness:**
 - Simplifying features can reduce computation costs but may make the model less robust to variations in the input data.
- **Latency vs. Real-Time Needs:**
 - Optimizing for low latency might require sacrificing some accuracy to meet real-time performance requirements on constrained hardware.

By iteratively testing and validating these trade-offs, the final model can strike a practical balance between fast inference and acceptable accuracy, suitable for the target application.

Experiments

EON Tuner

Impulse design

Create impulse

MFCC

Classifier

Retrain model

Live classification

Model testing

Perf. calibration

Deployment

Upgrade Plan

Get access to higher job limits, collaborators and a full commercial license.

View plans

Impulse #1

An impulse takes raw data, uses signal processing to extract features, and then uses a learning block to classify new data.

Time series data

Input axes
audio

Window size
1,000 ms.

Window increase (stride)
500 ms.

Frequency (Hz)
16000

Zero-pad data

Audio (MFCC)

Name
MFCC

Input axes (1)
Signal
audio

Classification

Name
Classifier

Input features
☒ MFCC

Output features
3 (helloworld, noise, unknown)

Output features

3 (helloworld, noise, unknown)

Save Impulse

EDGE IMPULSE

Dashboard

Devices

Data acquisition

Experiments

EON Tuner

Impulse design

Create impulse

MFCC

Classifier

Upgrade Plan

Get access to higher job limits, collaborators and a full commercial license.

View plans

Jagadeesh Kumar P / Keyword-spotting PERSONAL

Target: Arduino Nano 33 ...

Parameters

Generate features

Raw data

Show: All labels

helloworld.1ncrfnai.s1 (helloworld)

0ms 667ms 1335ms 2002ms 2670ms 3337ms 4005ms 4672ms 5340ms 6007ms 6675ms 7342ms 8010ms 8677ms 9345ms

0:00 / 0:01

Raw features

Label

0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...

helloworld

Parameters

Mel Frequency Cepstral Coefficients

Autotune parameters

DSP result

Cepstral Coefficients

10 9 8 7 6 5 4 3 2

10 9 8 7 6 5 4 3 2

EDGE IMPULSE

- Dashboard
- Devices
- Data acquisition
- Experiments
- EON Tuner
- Impulse design
 - Create impulse
 - MFCC
 - Classifier

Upgrade Plan

Get access to higher job limits, collaborators and a full

Parameters

Autotune parameters

Mel Frequency Cepstral Coefficients

Number of coefficients ②

Frame length ②

Frame stride ②

Filter number ②


FFT length ②

Normalization window size ②

Low frequency ②

High frequency ②

Pre-emphasis



Processed features

0.1901, 0.1066, 2.0978, -0.2058, 1.0188, 1.0947, 0.2244, 0.8897, 0.3878, 0.3037,...

On-device performance ②

PROCESSING TIME
258 ms.

PEAK RAM USAGE
15 KB

EDGE IMPULSE

- Dashboard
- Devices
- Data acquisition
- Experiments
- EON Tuner
- Impulse design
 - Create impulse
 - MFCC
 - Classifier

Upgrade Plan

Get access to higher job limits, collaborators and a full commercial license.

View plans

Neural Network settings

Training settings

Number of training cycles ②

Use learned optimizer ② ☐

Learning rate ②

Training processor ②

Advanced training settings

Audio training options

Data augmentation ② ☐

Neural network architecture

Architecture presets ② 1D Convolutional (Default) 2D Convolutional

Input layer (650 features)

Training output

Model

Model version: ② Quantized (int8)

Last training performance (validation set)

ACCURACY
88.5%

LOSS
0.27

Confusion matrix (validation set)

	HELLOWORLD	NOISE	UNKNOWN
HELLOWORLD	92.2%	6.3%	1.6%
NOISE	18.8%	77.7%	3.6%
UNKNOWN	9.8%	3.3%	87.0%
F1 SCORE	0.92	0.77	0.89

Metrics (validation set)

METRIC	VALUE
Area under ROC Curve ②	0.97
Weighted average Precision ②	0.89
Weighted average Recall ②	0.89

EDGE IMPULSE

- Dashboard
- Devices
- Data acquisition
- Experiments
- EON Tuner
- Impulse design
 - Create impulse
 - MFCC
 - Classifier

Upgrade Plan

Get access to higher job limits, collaborators and a full commercial license.

View plans

Summary

Model version: ② Unoptimized (float32)

Name

Label

CATEGORY	COUNT
helloworld	0
noise	0
unknown	0
uncertain	1

Detailed result

Show only unknowns

TIMESTAMP	HELLOWORLD	NOISE	UNKNOWN
0	0.29	0.29	0.42

RAW DATA


unknown.28e47b1a_nohash_3.wav.1ncrnk0h



Raw features

-134, -178, -225, -247, -225, -173, -180, -50, -5, 44, 70, 74, 62, 62, 59, 30, 2...

MFCC



EDGE IMPULSE

Dashboard

Devices

Data acquisition

Experiments

EON Tuner

Impulse design

Create impulse

MFCC

Classifier

Upgrade Plan

Get access to higher job limits, collaborators and a full commercial license.

A binary containing both the Edge Impulse data acquisition client and your full impulse.

MODEL OPTIMIZATIONS

Model optimizations can increase on-device performance but may reduce accuracy.



EON™ Compiler

Same accuracy, 40% less RAM, 49% less ROM.

Quantized (int8)

Selected ✓

	MFCC	CLASSIFIER	TOTAL
LATENCY	258 ms.	4 ms.	262 ms.
RAM	15.4K	3.8K	15.4K
FLASH	-	31.9K	-
ACCURACY			81.77%

Unoptimized (float32)

Select

	MFCC	CLASSIFIER	TOTAL
LATENCY	258 ms.	205 ms.	463 ms.
RAM	15.4K	7.0K	15.4K
FLASH	-	28.1K	-
ACCURACY			82.03%

Estimate for Arduino Nano 33 BLE Sense (Cortex-M4F 64MHz).

EDGE IMPULSE

Dashboard

Devices

Data acquisition

Experiments

EON Tuner

Impulse design

Create impulse

MFCC

Classifier

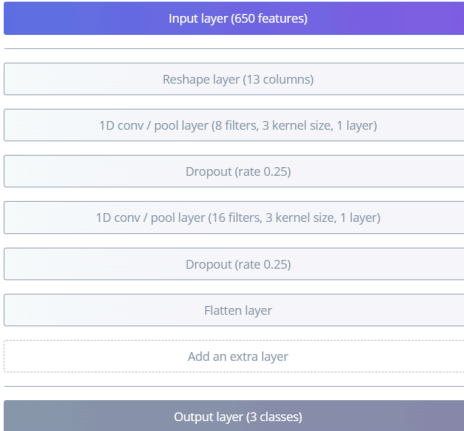
Upgrade Plan

Get access to higher job limits, collaborators and a full commercial license.

View plans

Neural network architecture

Architecture presets 1D Convolutional (Default) 2D Convolutional




Save & train

Metrics (validation set)

METRIC	VALUE
Area under ROC Curve	0.97
Weighted average Precision	0.89
Weighted average Recall	0.89
Weighted average F1 score	0.89

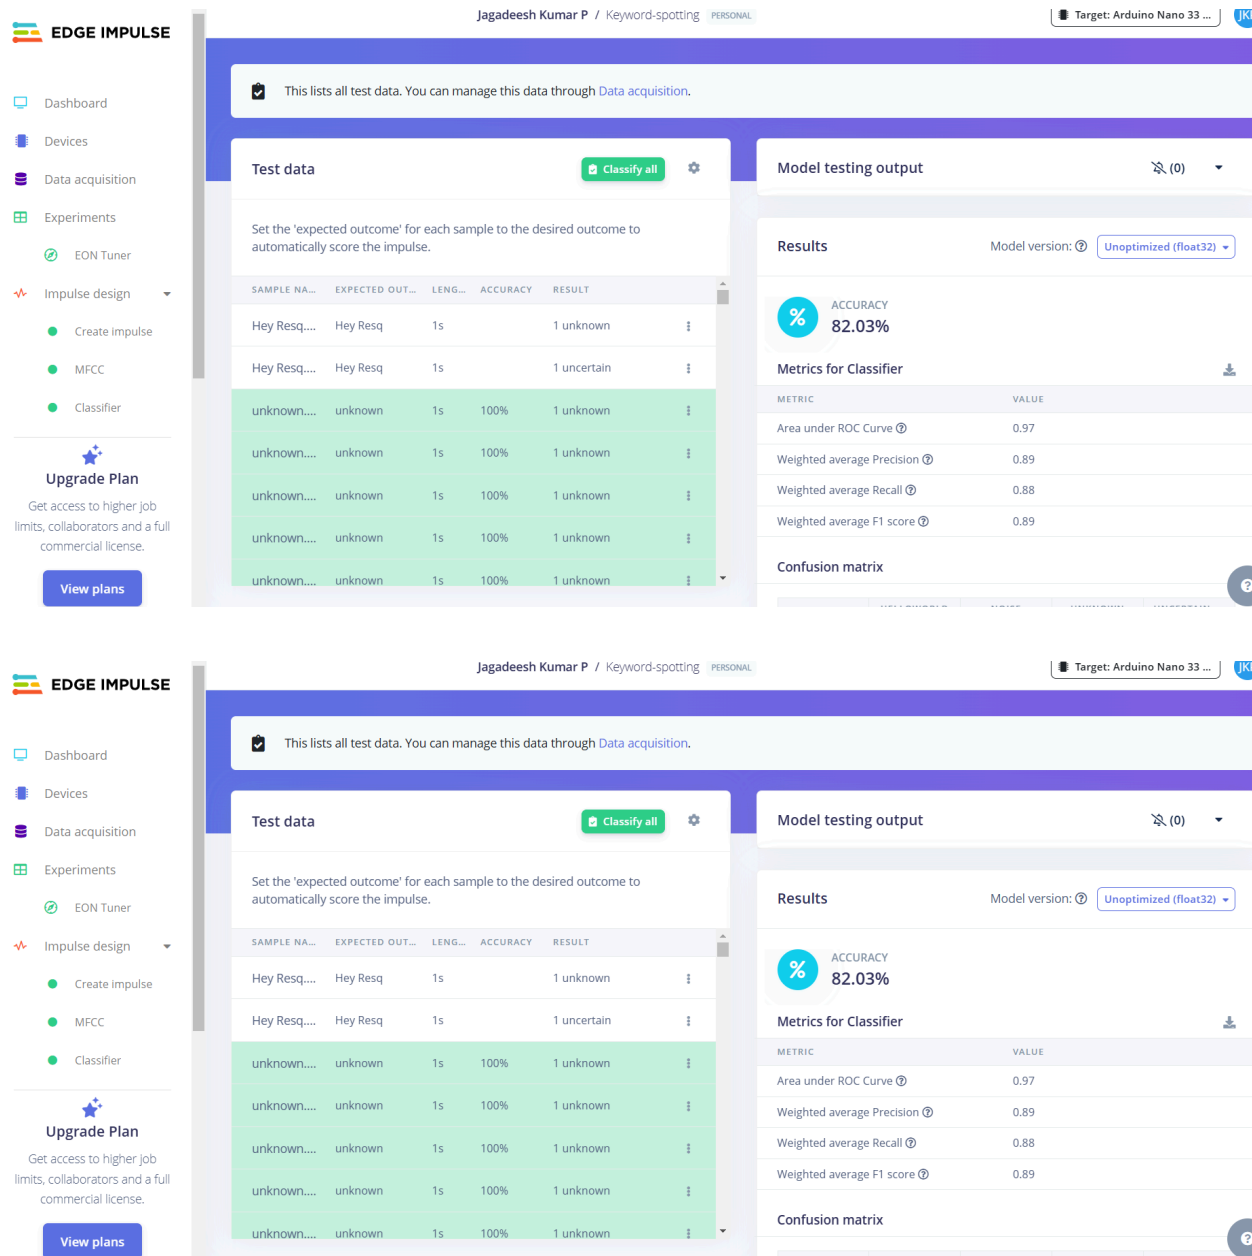
Data explorer (full training set)



On-device performance

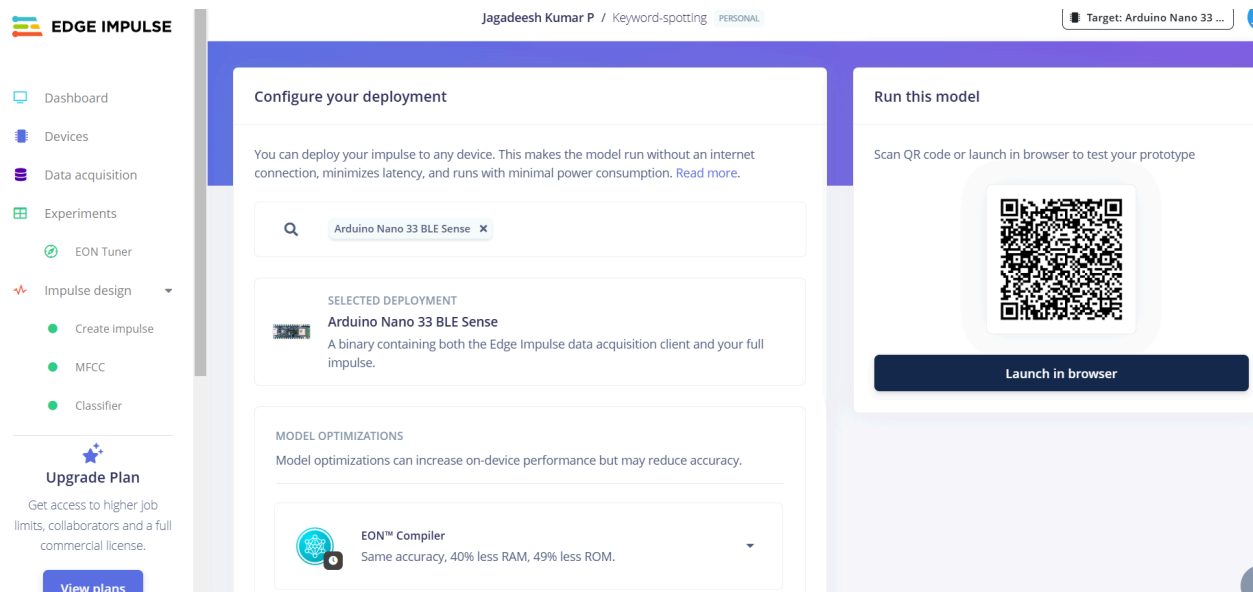
Engine: EON™ Compiler

INFERRING ...	PEAK RAM USA...	FLASH USAGE
4 ms.	3.8K	31.9K



Deployment to Smartphone:

- Process:** The deployment process to the smartphone was fairly straightforward, as Edge Impulse provides an easy-to-use interface to export models to mobile apps. The model was integrated into an app that communicates with the phone's microphone to perform real-time inference.
- Observations:** The accuracy on the smartphone was good, but not as high as the performance on the Arduino Nano 33 BLE Sense. This might be due to various factors like app limitations, processing power, or the environment in which the test was conducted (e.g., noise, microphone quality).



2. Videos:

- Record and provide links to:
 - The keyword-spotting model working on your smartphone.
 - The keyword-spotting model working on the embedded Arduino board.

Videos Uploaded

3. Reflections:

- Share your experience deploying the model to your smartphone and Arduino board. Mention any technical difficulties or interesting observations.

Deployment to Smartphone:

- **Process:** The deployment process to the smartphone was fairly straightforward, as Edge Impulse provides an easy-to-use interface to export models to mobile apps. The model was integrated into an app that communicates with the phone's microphone to perform real-time inference.
- **Observations:** The accuracy on the smartphone was good, but not as high as the performance on the Arduino Nano 33 BLE Sense. This might be due to various factors like app limitations, processing power, or the environment in which the test was conducted (e.g., noise, microphone quality).

Deployment to Arduino Nano 33 BLE Sense:

- **Process:** The deployment to the Arduino board involved converting the model into a format that the board could run (e.g., converting it to a TensorFlow Lite model). Edge Impulse's platform made the conversion process seamless, with easy integration into the Arduino environment.
- **Observations:** The Arduino Nano 33 BLE Sense performed better with the model in terms of accuracy. This was likely due to the simplicity of the hardware and the focus on low-power, low-latency inference. The BLE Sense board is designed with efficient sensor and signal processing capabilities, which could have contributed to this result.
- **Challenges:** The main difficulty was ensuring the model size was small enough to fit within the limited memory of the Arduino board. Optimizing the model to run efficiently on such a constrained device was crucial to achieving a good performance.
- **Interesting Points:** The BLE module's performance was surprisingly good considering its limited resources. This showed the power of TinyML and the importance of optimizing models for the target hardware.